

V2A2_Regression

December 14, 2018

```
In [1]: #!/usr/bin/env python
        # V2A2_Regression.py
        # Programmgeruest zu Versuch 2, Aufgabe 2

import numpy as np
import scipy.spatial
from random import randint

# -----
# base class for regressifiers
# -----
class Regressifier:
    """
    Abstract base class for regressifiers
    Inherit from this class to implement a concrete regression algorithm
    """

    def fit(self,X,T):          # train/compute regression with lists of feature vectors
        """
        Train regressifier by training data X, T, should be overwritten by any derived
        :param X: Data matrix of size NxD, contains in each row a data vector of size D
        :param T: Target vector matrix of size NxK, contains in each row a target vector of size K
        :returns: -
        """
        pass

    def predict(self,x):        # predict a target vector given the data vector x
        """
        Implementation of the regression algorithm; should be overwritten by any derived
        :param x: test data vector of size D
        :returns: predicted target vector
        """
        return None

    def crossvalidate(self,S,X,T,dist=lambda t: np.linalg.norm(t)): # do a S-fold cross validation
        """
        Do a S-fold cross validation
        """
```

```

:param S: Number of parts the data set is divided into
:param X: Data matrix (one data vector per row)
:param T: Matrix of target vectors; T[n] is target vector of X[n]
:param dist: a function dist(t) returning the length of vector t (default=Eukli
:returns (E_dist,sd_dist,E_min,E_max) : mean, standard deviation, minimum, and
:returns (Erel_dist,sdrel_dist,Erel_min,Erel_max) : mean, standard deviation,
"""

X,T=np.array(X),np.array(T)                                # ensure array type
N=len(X)                                                    # N=number of data vectors
perm = np.random.permutation(N)                             # do a random permutation
X1,T1=[X[i] for i in perm], [T[i] for i in perm]            # ... to get random partiti
idxS = [range(i*N//S,(i+1)*N//S) for i in range(S)]         # divide data set into S p
E_dist,E_dist2,E_max,E_min=0,0,-1,-1                        # initialize first two mom
Erel_dist,Erel_dist2,Erel_max,Erel_min=0,0,-1,-1            # initialize first two mom
for idxTest in idxS:                                       # loop over all possible t
    # (i) generate training and testing data sets and train classifier
    idxLearn = [i for i in range(N) if i not in idxTest]
    if(S<=1): idxLearn=idxTest
    X_learn, T_learn = np.array([X1[i] for i in idxLearn]), np.array([T1[i] for
    X_test, T_test = np.array([X1[i] for i in idxTest ]), np.array([T1[i] for
    self.fit(X_learn,T_learn)                               # train regressifier
    # (ii) test regressifier
    for i in range(len(X_test)): # loop over all data vectors to be tested
        # (ii.a) regress for i-th test vector
        xn_test = X_test[i].T                               # data vector for test
        t_test = self.predict(xn_test)                       # predict target value
        # (ii.b) check for regression errors
        t_true = T_test[i].T                                # true target value
        d=dist(t_test-t_true)                                # (Euklidean) distance
        dttrue=dist(t_true)                                  # length of t_true
        E_dist = E_dist+d                                    # sum up distances (for
        E_dist2 = E_dist2+d*d                                # sum up squared distan
        if(E_max<0)or(d>E_max): E_max=d                      # collect maximal error
        if(E_min<0)or(d<E_min): E_min=d                      # collect minimal error
        drel=d/dttrue
        Erel_dist = Erel_dist+drel                           # sum up relative dist
        Erel_dist2 = Erel_dist2+(drel*drel)                  # sum up squared relat
        if(Erel_max<0)or(drel>Erel_max): Erel_max=drel       # collect maximal rela
        if(Erel_min<0)or(drel<Erel_min): Erel_min=drel       # collect minimal rela
    E_dist = E_dist/float(N)                                  # estimate of first mo
    E_dist2 = E_dist2/float(N)                                # estimate of second m
    Var_dist = E_dist2-E_dist*E_dist                          # variance of error
    sd_dist = np.sqrt(Var_dist)                               # standard deviation o
    Erel_dist = Erel_dist/float(N)                             # estimate of first mo
    Erel_dist2 = Erel_dist2/float(N)                           # estimate of second m
    Varrel_dist = Erel_dist2-Erel_dist*Erel_dist              # variance of error
    sdrel_dist = np.sqrt(Varrel_dist)                         # standard deviation o
    return (E_dist,sd_dist,E_min,E_max), (Erel_dist,sdrel_dist,Erel_min,Erel_max)

```

and maximum error (f

```
In [2]: # -----
# DataScaler: scale data to standardize data distribution (for mean=0, standard deviat
# -----
class DataScaler:
    """
    Class for standardizing data vectors
    Some regression methods require standardizing of data before training to avoid num
    """

    def __init__(self,X):
        """
        Constructor: Set parameters (mean, std,...) to standardize data matrix X
        :param X: Data matrix of size NxD the standardization parameters (mean, std, .
        :returns: object of class DataScaler
        """
        self.meanX = np.mean(X,0)          # mean values for each feature column
        self.stdX = np.std(X,0)             # standard deviation for each feature column
        if isinstance(self.stdX,(list,tuple,np.ndarray)):
            self.stdX[self.stdX==0]=1.0 # do not scale data with zero std (that is, co
        else:
            if(self.stdX==0): self.stdX=1.0 # in case stdX is a scalar
        self.stdXinv = 1.0/self.stdX        # inverse standard deviation

    def scale(self,x):
        """
        scale data vector (or data matrix) x to mean=0 and s.d.=1
        :param x: data vector or data matrix
        :returns: scaled (standardized) data vector or data matrix
        """
        return np.multiply(x-self.meanX,self.stdXinv)

    def unscale(self,x):
        """
        unscale data vector (or data matrix) x to original data ranges
        :param x: standardized data vector or data matrix
        :returns: unscaled data vector or data matrix
        """
        return np.multiply(x,self.stdX)+self.meanX

    def printState(self):
        """
        print standardization parameters (mean value, standard deviation (std), and in
        """
        print("mean=",self.meanX, " std=",self.stdX, " std_inv=",self.stdXinv)

In [3]: from itertools import *
```

```

from functools import reduce
# -----
# function to compute polynomial basis functions
# -----
def phi_polynomial(x,deg=1):          # x should be list or np.array or 1xD matrix; r
    """
    polynomial basis function vector; may be used to transform a data vector x into a
    :param x: data vector to be transformed into a feature vector
    :param deg: degree of polynomial
    :returns phi: feature vector
    Example: phi_polynomial(x,3) returns for one-dimensional x the vector [1, x, x*x,
    """
    x=np.array(np.mat(x))[0]          # ensure that x is a 1D array (first row of x)
    D=len(x)
    #assert (D==1) or ((D>1) and (deg<=3)), "phi_polynomial(x,deg) not implemented for
    if(D==1):
        phi = np.array([x[0]**i for i in range(deg+1)])
    else:
        phi = np.array(phi_helper(x.tolist(),deg))
        #phi = np.array([])
        #if(deg>=0):
        #    phi = np.concatenate((phi,[1]))          # include degree 0 terms
        #    if(deg>=1):
        #        phi = np.concatenate((phi,x))          # includes degree 1 terms
        #        if(deg>=2):
        #            for i in range(D):
        #                phi = np.concatenate(( phi, [x[i]*x[j] for j in range(i+1)] ).
        #            if(deg>=3):
        #                for i in range(D):
        #                    for j in range(i+1):
        #                        phi = np.concatenate(( phi, [x[i]*x[j]*x[k] for k in
        #                # EXTEND CODE HERE FOR deg>3!!!!
    return phi.T # return basis function vector (=feature vector corresponding to dat

def phi_helper(x, deg):
    if deg <= 0:
        return [1]
    else:
        return phi_helper(x, deg-1) + [reduce(lambda a,b:a*b,combi) for combi in combi

In [4]: # -----
# Least Squares (ML) linear regression with sum of squares Regularization,
# -----
class LSRRRegressor(Regressor):
    """
    Class for Least Squares (or Maximum Likelihood) Linear Regressifier with sum of sq
    """

```

```

def __init__(self,lambda=0,phi=lambda x: phi_polynomial(x,1),flagSTD=0,eps=0.000001):
    """
    Constructor of class LSRegressifier
    :param lambda: Regularization coefficient lambda
    :param phi: Basis-functions used by the linear model (default linear polynomial)
    :param flagSTD: If >0 then standardize data X and target values T (to mean 0 and s.d.=1)
    :param eps: maximal residual value to tolerate (instead of zero) for numerical stability
    :returns: -
    """
    self.lambda=lambda          # set regression parameter (default 0)
    self.phi=phi                 # set basis functions used for linear regression (default linear polynomial)
    self.flagSTD=flagSTD;       # if flag >0 then data will be standardized, i.e., scaled to mean 0 and s.d.=1
    self.eps=eps;               # maximal residual value to tolerate (instead of zero)

def fit(self,X,T,lambda=None,phi=None,flagSTD=None): # train/compute LS regression
    """
    Train regressifier (see lecture manuscript, theorem 3.11, p33)
    :param X: Data matrix of size NxD, contains in each row a data vector of size D
    :param T: Target vector matrix of size NxK, contains in each row a target vector of size K
    :param lambda: Regularization coefficient lambda
    :param phi: Basis-functions used by the linear model (default linear polynomial)
    :param flagSTD: If >0 then standardize data X and target values T (to mean 0 and s.d.=1)
    :returns: flagOK: if >0 then all is ok, otherwise matrix inversion was bad conditioned
    """
    # (i) set parameters
    if lambda==None: lambda=self.lambda          # reset regularization coefficient?
    if phi==None: phi=self.phi                   # reset basis functions?
    if flagSTD==None: flagSTD=self.flagSTD        # standardize data vectors?
    # (ii) scale data for mean=0 and s.d.=1 ?
    if flagSTD>0:
        self.datascalerX=DataScaler(X)          # create datascaler for data matrix X
        self.datascalerT=DataScaler(T)          # create datascaler for target matrix T
        X=self.datascalerX.scale(X)              # scale all features (=columns) of data
        T=self.datascalerT.scale(T)              # ditto for target matrix T
    # (iii) compute weight matrix and check numerical condition
    flagOK,maxZ=1,0;                             # if <1 then matrix inversion is numerically bad
    try:
        self.N,self.D = X.shape                 # data matrix X has size N x D (N is number of samples, D is number of features)
        self.M = self.phi(self.D*[0]).size       # get number of basis functions
        #self.K = T.shape[1]                     # DELTE dummy code (just required for some tests)
        PHI = np.array([np.transpose(phi(x)) for x in X]) # compute basis functions for all data vectors
        PHIT_PHI_lambdaI = np.add(np.dot(PHI.T,PHI),np.dot(self.lambda,np.identity(self.M))) # compute PHIT_PHI_lambdaI
        #print(PHIT_PHI_lambdaI)
        # REPLACE dummy code: compute PHIT_PHI_lambdaI_inv
        PHIT_PHI_lambdaI_inv = np.linalg.pinv(PHIT_PHI_lambdaI) # REPLACE dummy code
        self.W_LSR = np.dot(np.dot(PHIT_PHI_lambdaI_inv,PHI.T),T) # REPLACE dummy code
    except:
        flagOK=0
    # (iv) check numerical condition

```

```

        Z=np.dot(PHIT_PHI_lambdaI,PHIT_PHI_lambdaI_inv)-np.identity(self.M) # REPLACE dummy code: Compute maximum (absolute) component
        maxZ = np.max(Z) # REPLACE dummy code: Compute maximum (absolute) component
        assert maxZ<=self.eps # maxZ should be <eps for good condition
    except:
        flagOK=0;
        print("EXCEPTION DUE TO BAD CONDITION:flagOK=", flagOK, " maxZ=", maxZ, " ")
        raise
    return flagOK

def predict(self,x,flagSTD=None): # predict a target value given data vector x
    """
    predicts the target value y(x) for a test vector x
    :param x: test data vector of size D
    :param flagSTD: If >0 then standardize data X and target values T (to mean 0 and std 1)
    :returns: predicted target vector y of size K
    """
    if flagSTD==None: flagSTD=self.flagSTD # standardization?
    if flagSTD>0: x=self.datascalerX.scale(x) # if yes, then scale x before computing phi_of_x
    phi_of_x = self.phi(x) # compute feature vector phi_of_x
    y=np.dot(self.W_LSR.T, phi_of_x) # REPLACE dummy code: compute prediction
    if flagSTD>0: y=self.datascalerT.unscale(y) # scale prediction back to original scale
    return y # return prediction y for data vector x

```

```

In [5]: # -----
# KNN regression
# -----

class KNNRegressifier(Regressifier):
    """
    Class for fast K-Nearest-Neighbor-Regression using KD-trees
    """

    def __init__(self,K,flagKLinReg=0):
        """
        Constructor of class KNNRegressifier
        :param K: number of nearest neighbors that are used to compute prediction
        :param flagKLinReg: if >0 then the do a linear (least squares) regression on the the
                           otherwise just take the mean of the K nearest neighbors target values
        :returns: -
        """
        self.K = K # K is number of nearest-neighbors
        self.X, self.T = [],[] # initially no data is stored
        self.flagKLinReg=flagKLinReg # if flag is set then do a linear regression

    def fit(self,X,T): # train/compute regression with lists of data vectors X and target values T
        """
        Train regressifier by storing X and T and by creating a KD-Tree based on X
        :param X: Data matrix of size NxD, contains in each row a data vector of size D
        :param T: Target vector matrix of size NxK, contains in each row a target vector of size K
        """

```

```

        :returns: -
        """
        self.X, self.T = np.array(X), np.array(T)    # just store feature vectors X and T
        self.N, self.D = self.X.shape                # store data number N and dimension D
        self.kdtree = scipy.spatial.KDTree(self.X)    # do an indexing of the feature vectors

def predict(self, x, K=None, flagKLinReg=None):      # predict a target value given data x
    """
    predicts the target value y(x) for a test vector x
    :param x: test data vector of size D
    :param K: number of nearest neighbors that are used to compute prediction
    :flagKLinReg: if >0 then the do a linear (least squares) regression on the the
                  otherwise just take the mean of the K nearest neighbors target values
    :returns: predicted target vector of size K
    """
    if(K==None): K=self.K                            # do a K-NN search...
    if(flagKLinReg==None): flagKLinReg=self.flagKLinReg # if flag >0 then do a regression
    nn = self.kdtree.query(x, K)                      # get indexes of K nearest neighbors
    if K==1: idxNN=[nn[1]]                            # cast nearest neighbor indexes nn to int
    else: idxNN=nn[1]
    t_out=0
    if(self.flagKLinReg==0):
        # just take mean value of KNNs
        t_out=np.mean([self.T[i] for i in idxNN])
    else:
        # do a linear regression of the KNNs
        lsr=LSRRegressor(lmbda=0.0001, phi=lambda x: phi_polynomial(x, 1), flagSTD=0)
        lsr.fit(self.X[idxNN], self.T[idxNN])
        t_out=lsr.predict(x)
    return t_out

```

```

In [17]: # *****
# __main__
# Module test
# *****

```

```

if __name__ == '__main__':
    print("\n-----")
    print("Example: 1D-linear regression problem")
    print("-----")
    # (i) generate data
    N=100
    w0, w1=4, 2                                # parameters of line
    X=np.zeros((N, 1))                          # x data: allocate Nx1 matrix as numpy ndarray
    X[:, 0]=np.arange(0, 50.0, 50.0/N)          # equidistant sampling of the interval [0, 50)
    T=np.zeros((N, 1))                          # target values: allocate Nx1 matrix as numpy ndarray
    sd_noise = 1.0                             # noise power (=standard deviation)
    T=T+w1*X+w0 + np.random.normal(0, sd_noise, T.shape) # generate noisy target values

```

```

par_lambda = 0                # regularization parameter
print("X=",X)
print("T=",T)

# (ii) define basis functions (phi should return list of basis functions; x should
deg=2;                          # degree of polynomial
phi=lambda x: phi_polynomial(x,1) # define phi by polynomial basis-functions up
print("phi(4)=", phi([4]))        # print basis function vector [1, x, x*x ..
print("phi([1,2])=", phi([1,2]))  # print basis function vector for two-dim.

# (iii) compute LSR regression
print("\n-----")
print("Do a Least-Squares-Regression")
print("-----")
lmbda=0;
lsr = LSRRegressifier(lmbda,phi)
lsr.fit(X,T)
print("lsr.W_LSR=",lsr.W_LSR)      # weight vector (should be approximately
x=np.array([3.1415]).T
print("prediction of x=",x,"is y=",lsr.predict(x))

# do S-fold crossvalidation
S=3
err_abs,err_rel = lsr.crossvalidate(S,X,T)
print("LSRRegression cross-validation: absolute errors (E,sd,min,max)=", err_abs,

# (iv) compute KNN-regression
print("\n-----")
print("Do a KNN-Regression")
print("-----")
K=2;
knnr = KNNRegressifier(K)
knnr.fit(X,T)
print("prediction of x=",x,"is y=",knnr.predict(x))

# do S-fold crossvalidation
err_abs,err_rel = knnr.crossvalidate(S,X,T)
print("KNNRegression cross-validation: absolute errors (E,sd,min,max)=", err_abs,

```

Example: 1D-linear regression problem

```

X= [[ 0. ]
     [ 0.5]
     [ 1. ]
     [ 1.5]
     [ 2. ]

```


[2.5]
[3.]
[3.5]
[4.]
[4.5]
[5.]
[5.5]
[6.]
[6.5]
[7.]
[7.5]
[8.]
[8.5]
[9.]
[9.5]
[10.]
[10.5]
[11.]
[11.5]
[12.]
[12.5]
[13.]
[13.5]
[14.]
[14.5]
[15.]
[15.5]
[16.]
[16.5]
[17.]
[17.5]
[18.]
[18.5]
[19.]
[19.5]
[20.]
[20.5]
[21.]
[21.5]
[22.]
[22.5]
[23.]
[23.5]
[24.]
[24.5]
[25.]
[25.5]
[26.]

[26.5]
[27.]
[27.5]
[28.]
[28.5]
[29.]
[29.5]
[30.]
[30.5]
[31.]
[31.5]
[32.]
[32.5]
[33.]
[33.5]
[34.]
[34.5]
[35.]
[35.5]
[36.]
[36.5]
[37.]
[37.5]
[38.]
[38.5]
[39.]
[39.5]
[40.]
[40.5]
[41.]
[41.5]
[42.]
[42.5]
[43.]
[43.5]
[44.]
[44.5]
[45.]
[45.5]
[46.]
[46.5]
[47.]
[47.5]
[48.]
[48.5]
[49.]
[49.5]]
T= [[3.96539552]

[5.40900355]
[6.80479406]
[8.03458322]
[8.3053615]
[8.73249338]
[8.77151182]
[9.63643954]
[12.01447684]
[14.10503997]
[14.53639261]
[16.56342681]
[15.55803052]
[17.54062696]
[18.10718482]
[20.19849118]
[19.69947678]
[21.30766924]
[21.40690018]
[22.14653271]
[22.94023538]
[25.36725429]
[26.22817455]
[27.78178612]
[28.41925754]
[27.71492035]
[29.23188689]
[29.9793288]
[31.95354147]
[32.47758882]
[32.27269297]
[36.15945205]
[35.22589939]
[38.2674975]
[39.33027519]
[38.71212506]
[40.19112542]
[40.37822356]
[42.08476023]
[41.09623498]
[43.07393777]
[44.58435281]
[46.23278059]
[46.80294075]
[48.494492]
[47.73994133]
[50.94662245]
[50.73035166]
[50.82361583]

[54.41661601]
[54.32981823]
[55.41985461]
[54.82862886]
[58.04696979]
[57.30488021]
[59.91427266]
[60.44129026]
[62.53071307]
[60.97598519]
[62.26103967]
[64.32530732]
[64.36329645]
[68.68025287]
[66.79291027]
[68.05828692]
[69.51147636]
[69.78195668]
[72.01666756]
[72.378812]
[73.80618286]
[73.32671575]
[74.58556612]
[75.35734468]
[77.2098026]
[79.00187307]
[79.67579371]
[79.86904299]
[80.98003463]
[82.03965717]
[82.55033359]
[81.91388516]
[86.1249457]
[84.73057236]
[86.42908599]
[86.79082706]
[88.72028873]
[90.40730705]
[91.78703418]
[91.15956868]
[93.33433493]
[93.12277541]
[96.2559933]
[95.55246418]
[95.90050568]
[97.07683317]
[98.35082574]
[99.51604217]

```

[100.21538238]
[103.11867279]
[102.75984867]]
phi(4)= [1 4]
phi([1,2])= [1 1 2]

-----
Do a Least-Squares-Regression
-----
lsr.W_LSR= [[4.09052788]
[1.99478017]]
prediction of x= [3.1415] is y= [10.35712978]
LSRRegression cross-validation: absolute errors (E,sd,min,max)= (0.7210491003592423, 0.5070165

-----
Do a KNN-Regression
-----
prediction of x= [3.1415] is y= 9.203975679168057
KNNRegression cross-validation: absolute errors (E,sd,min,max)= (1.1743092886650586, 0.9170722

In [16]: temp = []
         for i,e in enumerate(KNNErrors):
             temp.append((e[2][3],i))
         temp.sort()
         #print(temp)

         smallest_mean= KNNErrors[1][2]
         smallest_sd= KNNErrors[1][2]
         smallest_min= KNNErrors[13][2]
         smallest_max= KNNErrors[1][2]
         print(smallest_mean)
         print(smallest_sd)
         print(smallest_min)
         print(smallest_max)
         print(KNNErrors[1])

(0.04137596473341779, 0.07148122842087293, 6.810069605730603e-05, 0.44728370446074633)
(0.04137596473341779, 0.07148122842087293, 6.810069605730603e-05, 0.44728370446074633)
(0.1297162287702424, 0.39638164915995594, 5.395317915671042e-05, 3.1265953272944778)
(0.04137596473341779, 0.07148122842087293, 6.810069605730603e-05, 0.44728370446074633)
(2, (1.0953802187892574, 0.8689819772601937, 0.003291802811396849, 3.884868724720441), (0.0413

```

0.1 a) Versuchen Sie zunächst den Aufbau des Moduls V2A2_Regression.py zu verstehen:

Betrachten Sie den Aufbau des Moduls durch Eingabe von `pydoc V2A2_Regressifier`. Welche Klassen gehören zu dem Modul und welchen Zweck haben sie jeweils?

- DataScaler: Standartisiert Daten
- Regressifier: Abstrakte Klasse für Regressierer
- KNNRegressifier: ermöglicht Regression mithilfe von Fast-KNN
- LSRRegressifier: ermöglicht Regression mithilfe von Fehlerquadratsummen

Betrachten Sie nun die Basis-Klasse Regressifier im Quelltext: Wozu dienen jeweils die Methoden `fit(self,X,T)`, `predict(self,x)` und `crossvalidate(self,S,X,T)` ?

- `fit(self,X,T)`: Verlangt nach der Implementierung einer Methode die den Regressierer mit X und T trainiert.
- `predict(self,x)`: Verlangt nach der Implementierung einer Methode die auf x eine Regression anwendet und eine Vorhersage zurückliefert
- `crossvalidate(self,S,X,T)`: Macht S-fache Kreuzvalidierung mit den Daten, Labels (X,T) und liefert als Ergebnis Informationen über die relativen und absoluten Fehlerwerte

Worin unterscheidet sich `crossvalidate(.)` von der entsprechenden Methode für Klassifikation (siehe vorigen Versuch)?

- hier kann man noch angeben mit welcher Längenfunktion gearbeitet werden soll, und es werden andere Fehlerstatistiken zurückgegeben

b) Betrachten Sie nun die Funktion `phi_polynomial(x,deg)`:

Was berechnet die Funktion? Welches Ergebnis liefert `phi_polynomial([3],5)`? Welches Ergebnis liefert `phi_polynomial([3,5],2)`?

```
In [74]: phi_polynomial([3],5)
```

```
Out[74]: array([ 1,  3,  9, 27, 81, 243])
```

```
In [75]: phi_polynomial([3,5],2)
```

```
Out[75]: array([ 1,  3,  5,  9, 15, 25])
```

- `phi_polynomial(x,deg)`: berechnet die Basisfunktion für die Werte x mit dem Grad deg

Geben Sie eine allgemeine Formel an für das Ergebnis von `phi_polynomial([x1,x2],2)`?

- $[1, x_1, x_2, (x_1)^2, x_1 \cdot x_2, (x_2)^2]$

Wozu braucht man diese Funktion im Zusammenhang mit Regression?

- Das Kreuzprodukt aus dieser Funktion und dem Zielwertvektor geben die Prognose bei der Regression

Bis zu welchem Polynomgrad kann die Funktion Basisfunktionen berechnen? Erweitern Sie die Funktion mindestens bis Grad 5.

- momentan bis Grad 3, siehe Verbesserung im Code

```
In [83]: phi_polynomial([3,5],5)
```

```
Out[83]: array([ 1, 3, 5, 9, 15, 25, 27, 45, 75, 125, 81,
               135, 225, 375, 625, 243, 405, 675, 1125, 1875, 3125])
```

c) Betrachten Sie die Klasse LSRRegressif:

Welche Art von Regressions-Modell berechnet diese Klasse?

- Fehlerquadratsummenregression mit Regularisierung

Wozu dienen jeweils die Parameter λ , ϕ , flagSTD und ϵ ?

- λ : Regularisierungsparameter
- ϕ : Basisfunktion
- flagSTD : gibt an ob die Daten Standardisiert werden sollen
- ϵ : höchster erlaubter Fehlerrestwert

Welche Rolle spielt hier die Klasse DataScaler? In welchen Methoden und zu welchem Zweck werden die Daten ggf. umskaliert? Welches Problem kann auftreten wenn man dies nicht tut? Wozu braucht man die Variablen Z und $\max Z$ in der Methode $\text{fit}()$?

- DataScaler wird in $\text{fit}()$ und $\text{predict}()$ verwendet um Daten zu Standardisieren
- nicht skalierte Daten werden können kleinste Abweichungen in Gewichten zu großen und/oder unterschiedlichen Änderungen führen.
- Z ist die Matrix mit den Fehlerwerten die durch Invertierung entstehen.
- $\max Z$ ist der grösste Wert aus Z und wird verwendet um zu bestimmen ob der Fehlerwert annehmbar ist

Vervollständigen Sie die Methoden $\text{fit}(\text{self}, X, T, \dots)$ und $\text{predict}(\text{self}, x, \dots)$ (vgl. vorige Aufgabe).

- siehe code

d) Betrachten Sie die Klasse KNNRegressif:

Welche Art von Regressions-Modell berechnet diese Klasse?

- fast K-Nearest-Neighbor-Regression

Wozu dienen jeweils die Parameter K und flagKLinReg ?

- K : Anzahl der NN die verwendet werden um eine Vorhersage zu treffen
- flagKLinReg : gibt an ob eine Regression verwendet werden soll oder nur der Mittelwert aus den NN genommen wird

Beschreiben Sie kurz in eigenen Worten (2-3 Sätze) auf welche Weise die Prädiktion $y(x)$ berechnet wird.

- Es werden die KNNs für x ermittelt. Damit wird ein LSRRegressif trainiert, der dann nach einer prediction gefragt wird.

e) Betrachten Sie abschließend den Modultest:

Beschreiben Sie kurz was im Modultest passiert.

- es wird ein Linearer Datensatz erstellt

- die Funktion phi wird neu definiert und auf Grad 2 festgesetzt
- LSR Regression wird durchgeführt und Kreuzvalidiert
- KNN Regression wird durchgeführt und Kreuzvalidiert

Welche Gewichte W werden gelernt? Wie lautet also die gelernte Prädiktionsfunktion? Welche Funktion sollte sich idealerweise (für N) ergeben?

$w_0 = [3.78705442]$
 $w_1 = [2.01036841]$
 $y = 3.787 + 2.010 \cdot x$
 $y = 4 + 2 \cdot x$ für N

Welche Ergebnisse liefert die Kreuzvalidierung? Was bedeuten die Werte?

- (E,sd,min,max)= (0.8335277318740052, 0.6423712068147653, 0.030450286307427632, 3.059745154084382)
- E ist der mittlere Fehlerwert
- sd ist die Standardabweichung der Fehler
- min ist der kleinste Fehlerwert
- max ist der größSte Fehlerwert

Vergleichen und Bewerten Sie die Ergebnisse von Least Squares Regression gegenüber der KNN-Regression (nach Optimierung der Hyper-Parameter , K, ...).

LSR: absolute errors (E,sd,min,max)= (0.7210491003592423, 0.5070165487566235, 0.0066007843635347285, 2.780939336314603)
KNN: absolute errors (E,sd,min,max)= (1.1743092886650586, 0.9170722237516987, 0.0038023455924474092, 4.2882465799825695)
LSR: relative errors (E,sd,min,max)= (0.02240702968734219, 0.028216719179408916, 0.0001525998169486947, 0.14635559861181296)
KNN: relative errors (E,sd,min,max)= (0.05354319692710508, 0.13202983166413382, 5.0979643788569236e-05, 1.0603171407681475)

- LSR scheint in fast allen Bereichen wesentlich bessere Ergebnisse als KNNR zu liefern.